

Mohammed Sedeg

Turkey (Open to relocate)

To the Mozn Team,

I am excited to apply for the Machine Learning Engineer role at Mozn. As an AI Engineer and Researcher with a deep personal and cultural connection to Saudi Arabia, I'm inspired by Mozn's leadership in advancing AI to achieve **Vision 2030**. The opportunity to help build scalable, real-world solutions in a region I care about deeply aligns perfectly with my mission as an engineer.

At **VisionCore**, I designed and deployed low-latency computer vision systems with sub-80ms edge inference on Jetson and Raspberry Pi. I optimized models using quantization, pruning, and clustering, and deployed YOLOv8 detectors while automating pipelines for labeling and anomaly detection. In parallel, as a freelance **LLM Engineer**, I fine-tuned models like GPT-2 and LLaMA using Hugging Face and LoRA, built RAG systems with LangChain and FAISS, and accelerated inference with quantization (gguf, ExLlama) and distributed training (PyTorch DDP).

My background strongly matches your key requirements:

- **Efficient model training and inference:** Leveraged DDP, AMP, and hardware-specific optimization.
- **Production deployment:** Experience with TFLite, ONNX, TorchServe, and Nvidia Triton.
- **Performance monitoring:** Familiar with Grafana, Prometheus, and deployment observability.
- **Infrastructure and automation:** Skilled in Docker, CI/CD, FastAPI, and AWS (including SageMaker).

As a native Arabic speaker with strong communication skills and a solid foundation in mechatronics and control systems, I bring both the technical depth and regional awareness needed to contribute meaningfully at Mozn.

I would be honored to join your mission of building world-class AI from the region, for the region—and beyond.

Warm regards,

Mohammed Sedeg
